

[Back](#)

**Your grade: 100%**

Your latest: **100%** • Your highest: **100%** • To pass you need at least 80%. We keep your highest score.

Next item →

## ^ Assignment Overview

Now, it's time to upload the results of building your own AI-powered podcast.

You will upload a course-related document (for example, the specialization narrative or selected readings) into NotebookLM. The tool will generate a personalized podcast or audio summary based on your document. Afterward, you'll reflect on three questions:

- How well did the tool summarize the material?
- What was included or left out?
- Did the output align with what you wanted?

Finally, you'll compare this experience with a simpler approach: asking a general AI system like ChatGPT to summarize “a Coursera course on generative AI” WITHOUT uploading the provided outline document. You'll most likely notice a response that is not aligned with this particular course.

This activity will help you understand the value of RAG-style tools while also showing their limitations. It will prepare you for deeper discussions in Week 4 on where AI can go wrong and why human judgment remains essential.



1. How well did the tool summarize the material? (Write at least 250-500 words)

1 / 1 point

The tool did a good job, but it kept things pretty surface-level. The summary correctly identified the big picture. It showed that the specialization moves learners through three stages: understanding AI basics, building practical skills, and eventually thinking at a systems level. It mentioned the right technical topics like RAG, multimodal tools, AI agents, bias, and hallucinations. And it ended on the right note, capturing the course's core message that people who use AI well will become harder to replace, not easier. So if someone had never seen the course before and read this summary, they would walk away with a reasonable idea of what it covers.

But it missed a lot of what makes the course special. The source document is not just a list of topics. It has a personality. There is a recurring metaphor of treating AI like a stranger on a bus, which shows up multiple times across different modules and gets more serious as the course goes on. The summary does not mention it at all. There is also a deliberate promise built into the course that you do not need math or coding experience to engage with the material, and that spirit is completely absent from the summary.

The course also has a very specific and memorable teaching approach around prompting. It dedicates an entire module to the idea, drawing on AI researcher Andrej Karpathy's concept of context engineering. The argument is that writing good prompts is just the tip of the iceberg, and that truly skilled AI users learn to manage the full context of a conversation. The summary reduces all of that to the phrase "effective prompting and context engineering," which is technically accurate but tells you almost nothing.

The ethical content also got flattened. The summary mentions bias and hallucinations, which is correct, but Module 4 goes much further. It covers data monetization, the environmental cost of running large models, real historical examples of AI failures like Microsoft's Tay chatbot, and the uncomfortable idea that when you use a free AI tool, your behavior is being studied and monetized. That is a much richer conversation than just "bias and hallucinations."

One of the most interesting parts of the course is where learners critique the AI's own evaluation of their work, developing what the course calls "meta-literacy." This is missing entirely from the summary.

The summary is accurate but thin. It covers the what but skips the how and the why. It reads like something written from the table of contents rather than from engaging with the course's ideas and teaching philosophy. For a quick one-paragraph overview, it works fine. But it would not give a prospective learner a real feel for what makes this course distinct, including the metaphors, the ethical depth, the specific way it teaches, and the consistent thread of personalized learning that runs through all three courses.



2. What was included or left out? (Write at least 250-500 words)

1 / 1 point

The summary got the big picture right. It correctly showed that the course moves learners through three stages: starting with AI basics, building practical skills, and finishing with more advanced thinking. This matches what the source document describes. It also mentioned the right technical topics. Things like RAG, multimodal tools, image and audio generation, and AI agents were all in the summary and are genuinely covered in the course. The ethical side was touched on. The summary noted that the program talks about responsible use and covers problems like bias and hallucinations. These are real topics in the course, especially in Module 4. The closing message was accurate too. The idea that learners will be able to work alongside AI and stay valuable in their careers matches the course's main motivational message, which appears at the start and end of the specialization.

A lot was missing, especially the parts that make the course feel unique. The "stranger on a bus" metaphor was never mentioned. This is one of the most recognizable parts of the course. It is introduced in Module 1 to encourage learners to approach AI in an open and curious way. It comes back in Module 3 to explain how earlier parts of a conversation still affect later responses. And it appears again in Module 4 with a more serious message about privacy and data use. Leaving it out means the summary missed one of the key threads that ties the whole course together.

The "no math or code required" promise was also left out. This is an important part of how the course describes itself to beginners. It tells learners who the course is for and helps set expectations. Someone considering the course would want to know this.

Vibe coding was not mentioned at all. The source describes it to show learners that they can use plain language to build programs. It is meant to make coding feel approachable for beginners while also being useful for people with more experience. It runs through Course 2 and connects to the broader learning theme.

The teaching approach around prompting was oversimplified. The course spends an entire module on prompting and context engineering, even referencing a well-known AI researcher named Andrej Karpathy. The summary reduced all of this to a short phrase that does not capture what is taught.

The ethical content was much thinner than what the source covers. Module 4 goes well beyond bias and hallucinations. It talks about how free AI tools use your behavior as data, the environmental cost of running large AI models, and real examples of AI going wrong like Microsoft's Tay chatbot. None of that made it into the summary.

The meta-literacy activity was completely left out. This is where learners get feedback on their work from an AI and then turn around and critique that feedback. The source treats this as one of the most important learning moments in the whole course. Leaving it out is a big gap.

The personalized learning thread was also missing. Throughout all three courses, the source uses AI powered personalized learning as a running example to make concepts feel real and relevant. It is not just a topic; it is part of how the course teaches. The summary did not reflect this at all.

The summary included the right topics at a broad level but left out most of what makes the course distinct. What was included was accurate. What was left out was often the most interesting and memorable parts of the curriculum.



3. Did the output align with what you wanted? (Write at least 250-500 words)

1 / 1 point

In some ways, yes. The summary got the big picture right. It showed that the course moves learners through three stages, starting with the basics, building up practical skills, and finishing with more advanced thinking. That matches what the course was designed to do. The ending message also came through well. The course puts a lot of effort into the idea that people who use AI well will become more valuable in their careers, not less. That idea showed up in the summary, and it is one of the most important points the course creators would want someone to walk away with. The summary also touched on ethics, noting that the course covers things like bias and hallucinations. The course does cover these topics, so that part was accurate.

In many ways though, the summary missed what makes the course feel different from other courses. The course was built with a very specific and friendly teaching style, and almost none of that came through. The course creators used a running metaphor of treating AI like a stranger on a bus to help learners feel comfortable and curious rather than intimidated. That metaphor shows up several times across different modules and gets more serious as the course goes on. The summary never mentioned it. The course also makes a clear promise to learners that they do not need any math or coding background to follow along. That promise is a big part of how the course welcomes beginners, and the summary left it out completely. The course also uses personalized learning as a running example throughout all three courses. It is not just a topic that comes up once. It is used again to show learners how AI can help them in their own lives and work. The summary did not reflect this at all. Vibe coding, which is the idea that anyone can use plain everyday language to build simple programs, was another big theme in the course that the summary skipped entirely. This was an important part of the course because it tried to make something that feels technical seem approachable and fun.

One of the most memorable parts of the course is an activity in Module 5 where learners get feedback on their work from an AI and then must judge whether that feedback was good or not. The course calls this meta-literacy, which just means the ability to think critically about the tools that are thinking about you. The summary never brought this up, even though the course treats it as one of the highlights. The ethics section also went much deeper in the source than the summary suggested. The course talked about how free AI tools collect and use your behavior as data, the amount of energy it takes to run large AI models, and real examples of AI going wrong in the past. The summary reduced all of this to just a few words. Overall, the summary was accurate at a basic level but felt thin. It covered the main topics but left out the personality, the teaching style, and the specific moments that make the course stand out. Someone reading only the summary would get a general sense of what the course covers but would have no real feel for what it is like to take it.



4. Compare this experience with a simpler approach: asking a general AI system like ChatGPT to summarize “a Coursera course on generative AI” from your own perspective (for example, “I’m a teacher focused on improving my skills in designing assignments” or “I’m a data analyst hoping to make my work more efficient”). Non-native speakers may also try generating the summary in their preferred language.

1 / 1 point

What did you notice from this comparison of general AI systems vs. RAG systems?

The biggest difference between the two systems was depth and specificity. The non-RAG system read the document and gave back an accurate but generic list of topics that could describe almost any AI course. It never mentioned the stranger on a bus metaphor, the vibe coding theme, or the meta-literacy activity, which are the details that make this course unique. The RAG system, on the other hand, had the actual source material to work with and was able to engage with the specific ideas, phrases, and activities in the document.

5. Now that you’ve created your AI-powered podcast, take a moment to reflect on the results. In the discussion forum, share your observations about what stood out to you. Consider questions like:

1 / 1 point

- What did your AI-generated podcast emphasize most strongly?

The podcast focused heavily on the journey from seeing AI as a simple prediction tool to understanding it as something that can plan and take actions on its own. It used the stranger on a bus metaphor to explain how learners should approach AI at the start, with curiosity but also caution. It also pushed the idea that AI is a force multiplier, meaning it helps you get more done rather than doing everything for you. And it kept coming back to the human and AI partnership message, especially the line that AI will not replace you if you are using AI to make yourself harder to replace.

Overall, the podcast did a good job of capturing the emotional arc of the course. It picked up the right metaphors, the right themes, and the right motivational message. What it emphasized most was the human side of the story, which makes sense because that is what the course leads with and closes with. The more technical details like specific module activities, the meta-literacy exercise, and the vibe coding theme were less present, but for a podcast format that is understandable. A podcast is meant to inspire and orient a listener, not cover every detail. So, the emphasis made sense for the format even if it left some of the richer course content on the table.



6. What surprised you about the results?

1 / 1 point

The biggest surprise is how well the podcast picked up on the emotional and motivational thread of the course. A lot of people assume AI tools are better at handling facts and lists than they are at capturing tone and feeling. But the podcast led with the human and AI partnership message and kept coming back to it, which is exactly what the course itself does. That is not an obvious thing to get right, and it is worth noting.

It is also surprising that the podcast organized its summary around three clear areas on its own. Nobody told it to structure the response that way. It read the material and decided that the shift from probability to agency, the practical metaphors, and the human AI partnership were the three most important threads. And honestly, that is a reasonable editorial decision. A human editor might have come up with something similar.

The stranger on a bus metaphor showing up is also notable. That metaphor is woven into the course across multiple modules and carries different meanings at different points. The fact that the podcast identified it as worth highlighting suggests the system was doing more than just pulling out keywords. It was picking up on what the course kept returning to.

What is less surprising is what the podcast left out. The more specific and technical details like the meta-literacy activity, the vibe coding theme, and the personalized learning thread did not make it in. That seems like a common pattern with AI generated summaries. They tend to gravitate toward the big ideas and skip the smaller but often more memorable moments.

The overall takeaway is that the podcast performed better than most people would expect on the qualitative and human centered parts of the course, and about as expected on the technical and specific parts.

^ Closing

As you compare your experience with others, you'll start to see that AI tools don't always make the same choices. This is part of their non-deterministic behavior, where outputs can vary depending on how the system processes information. By noticing these patterns, you'll begin developing the critical literacy needed to evaluate AI outputs thoughtfully and responsibly.